Clustering

* taken from Ron Yurko's slides

Q. Cluster NBA players into offensive Roles OR Archetypes based on How they play.

* First, need variables that describe How they play.

* Todd Whitehead: try using Synergy Play Type data!

```
[1] "FGA_freq_Spotup"           "FGA_freq_PRRollMan"          "FGA_freq_Postup"
[4] "FGA_freq_Cut"              "FGA_freq_OffRebound"         "FGA_freq_PRBallHandler"
[7] "FGA_freq_Isolation"        "FGA_freq_OffScreenOrHandoff"
```

* Then, with this data, how to cluster players??

What is **unsupervised learning?**

We have p variables for n observations $x_1, \ldots, x_n$, and for observation i:

$$x_{i1}, x_{i2}, \ldots, x_{ip}$$

- *unsupervised*: none of the variables are **response** variables, i.e., there are no labeled data

Think of unsupervised learning as **an extension of EDA...**

- $\Rightarrow$ **there is no unique right answer!**

What is clustering (aka cluster analysis)?

ISLR 10.3:

> *very broad set of techniques for finding subgroups, or clusters, in a dataset*

- observations **within** clusters are **more similar** to each other,
- observations **in different** clusters are **more different** from each other

How do we define **distance / dissimilarity** between observations?

- e.g. **Euclidean distance** between observations i and j

$$d(x_i, x_j) = \sqrt{(x_{i1} - x_{j1})^2 + \cdots + (x_{ip} - x_{jp})^2}$$

Units matter!

- one variable may *dominate* others when computing Euclidean distance because its range is much larger
- can standardize each variable / column of dataset to have mean 0 and standard divation 1 with `scale()`
- **but we may value the separation in that variable!** (so just be careful...)

What's the clustering objective?

- $C_1, \ldots, C_K$ are *sets* containing indices of observations in each of the K clusters
  - if observation i is in cluster k, then $i \in C_k$
- We want to minimize the **within-cluster variation** $W(C_k)$ for each cluster C_k and solve:

$$\underset{C_1,\ldots,C_K}{\text{minimize}} \left\{ \sum_{k=1}^{K} W(C_k) \right\}$$

- Can define using the **squared Euclidean distance** ($|C_k| = n_k =$ # observations in cluster k)

$$W(C_k) = \frac{1}{|C_k|} \sum_{i,j \in C_k} d(x_i, x_j)^2$$

- Commonly referred to as the within-cluster sum of squares (WSS)

Lloyd's algorithm

1) Choose K random centers, aka **centroids**

2) Assign each observation closest center (using Euclidean distance)

3) Repeat until cluster assignment stop changing:

- Compute new centroids as the averages of the updated groups
- Reassign each observations to closest center

Converges to a local optimum, not the global

Results will change from run to run (set the seed!)

Takes K as an input!

things to check when clustering

- do the units of the variables make sense?
- should you standardize the variables?
- is "nstart", the number of starting random configurations, large enough?

So, how do we choose the number of clusters?!

There is no universally accepted way to conclude that a particular choice of K is optimal!

Popular heuristic: elbow plot (use with caution)

Choose K where marginal improvements is low at the bend (hence the elbow)

This is just a guideline and should not dictate your choice of K!

Gap statistic is a popular choice (see `clusGap` function in `cluster` package)

Next Tuesday: model-based approach to choosing the number of clusters!

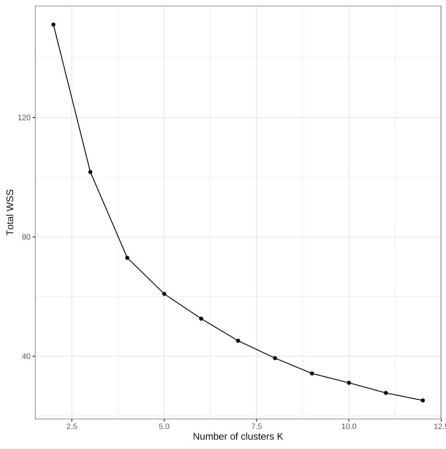

Better alternative to `nstart`: K-means++

Pick a random observation to be the center c_1 of the first cluster C_1

- This initializes a set $Centers = \{c_1\}$

Then for each remaining cluster $c^* \in 2, \ldots, K$:

- For each observation (that is not a center), compute $D(x_i) = \min_{c \in Centers} d(x_i, c)$
    - Distance between observation and its closest center $c \in Centers$
- Randomly pick a point x_i with probability: $p_i = \frac{D^2(x_i)}{\sum_{j=1}^{n} D^2(x_j)}$
- As distance to closest center increases $\Rightarrow$ probability of selection increases
- Call this randomly selected observation c^*, update $Centers = Centers \cup c^*$
    - Same as `centers = c(centers, c_new)`

Then run K-means using these $Centers$ as the starting points

Time to Cluster NBA players into offensive Roles by Synergy FGA play type frequencies

```
### Hard Clustering: K means
set.seed(387397)
km1 <- flexclust::kcca(mat_train, K1)
print(km1)
```

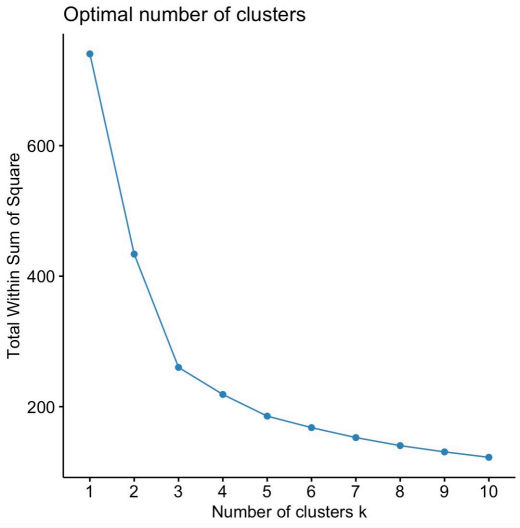

Elbow Plot
Began with K=3 clusters

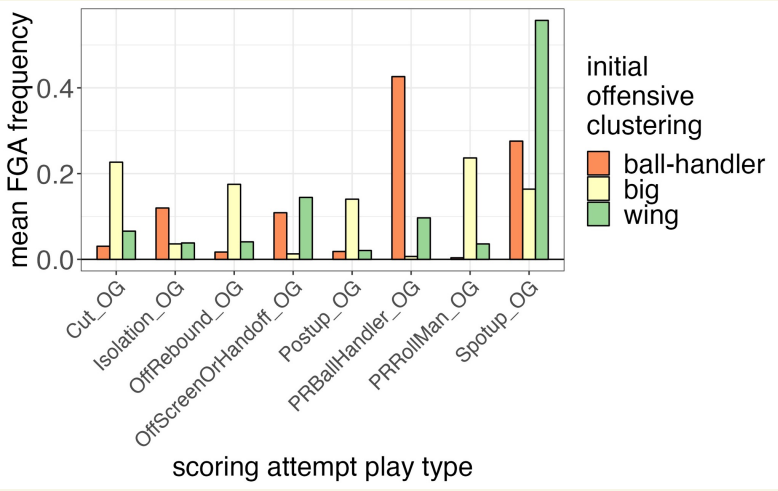

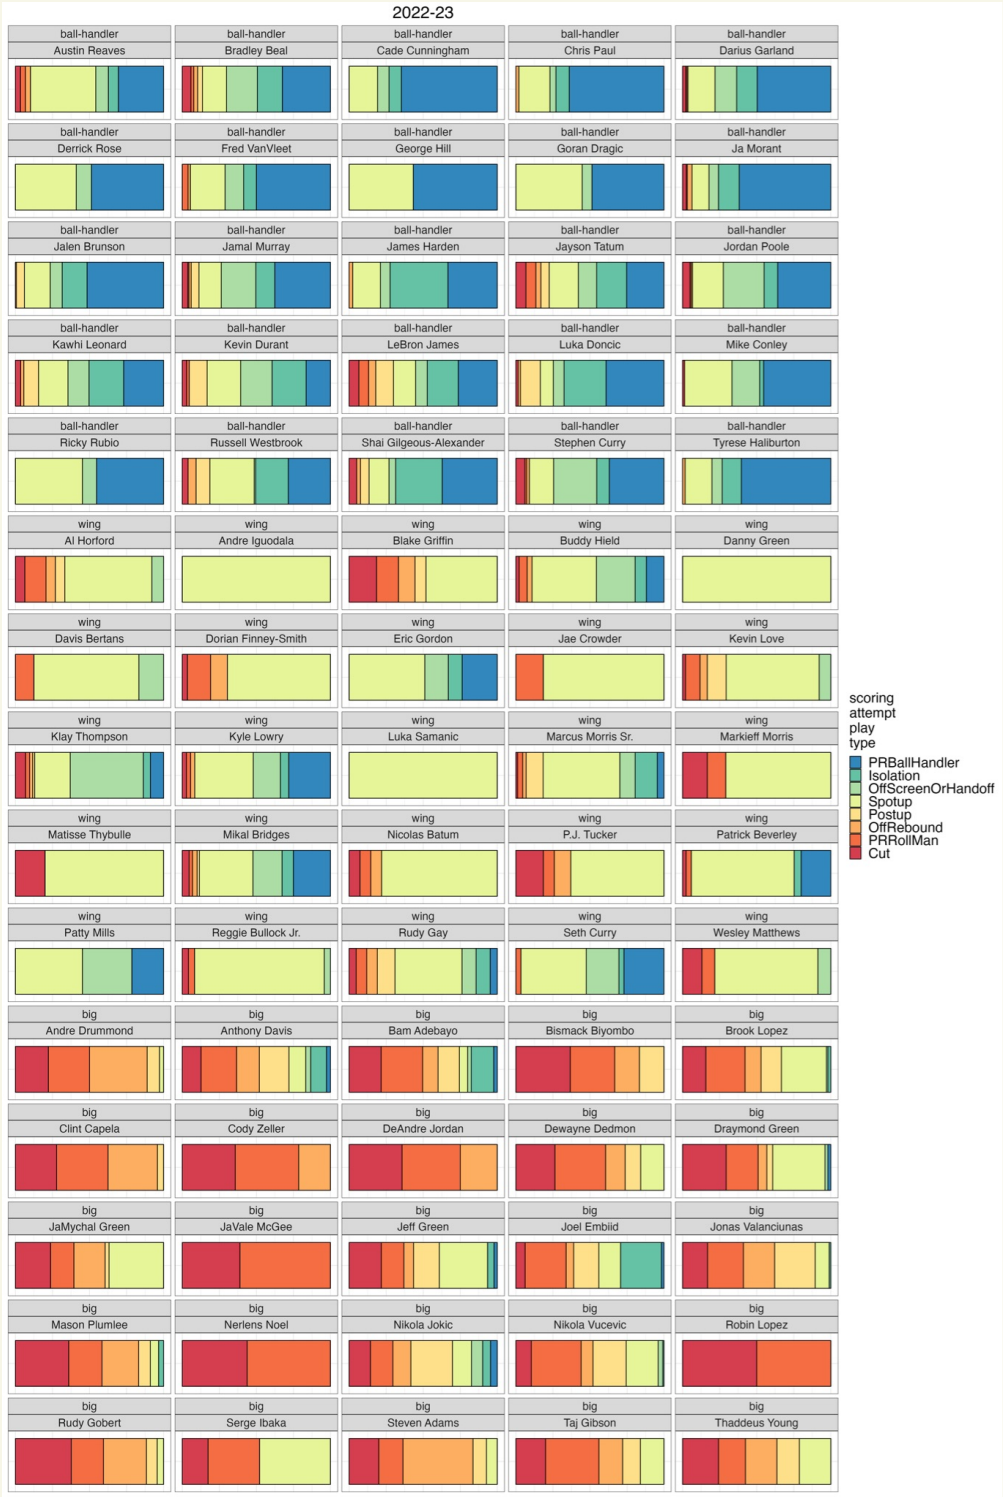
2022-23
ball-handler
Austin Reaves
ball-handler
Bradley Beal
ball-handler
Cade Cunningham
ball-handler
Chris Paul
ball-handler
Darius Garland
ball-handler
Derrick Rose
ball-handler
Fred VanVleet
ball-handler
George Hill
ball-handler
Goran Dragic
ball-handler
Ja Morant
ball-handler
Jalen Brunson
ball-handler
Jamal Murray
ball-handler
James Harden
ball-handler
Jayson Tatum
ball-handler
Jordan Poole
ball-handler
Kawhi Leonard
ball-handler
Kevin Durant
ball-handler
LeBron James
ball-handler
Luka Doncic
ball-handler
Mike Conley
ball-handler
Ricky Rubio
ball-handler
Russell Westbrook
ball-handler
Shai Gilgeous-Alexander
ball-handler
Stephen Curry
ball-handler
Tyrese Haliburton
wing
Al Horford
wing
Andre Iguodala
wing
Blake Griffin
wing
Buddy Hield
wing
Danny Green
wing
Davis Bertans
wing
Dorian Finney-Smith
wing
Eric Gordon
wing
Jae Crowder
wing
Kevin Love
wing
Klay Thompson
wing
Kyle Lowry
wing
Luka Samanic
wing
Marcus Morris Sr.
wing
Markieff Morris
wing
Matisse Thybulle
wing
Mikal Bridges
wing
Nicolas Batum
wing
P.J. Tucker
wing
Patrick Beverley
wing
Patty Mills
wing
Reggie Bullock Jr.
wing
Rudy Gay
wing
Seth Curry
wing
Wesley Matthews
big
Andre Drummond
big
Anthony Davis
big
Bam Adebayo
big
Bismack Biyombo
big
Brook Lopez
big
Clint Capela
big
Cody Zeller
big
DeAndre Jordan
big
Dewayne Dedmon
big
Draymond Green
big
JaMychal Green
big
JaVale McGee
big
Jeff Green
big
Joel Embiid
big
Jonas Valanciunas
big
Mason Plumlee
big
Nerlens Noel
big
Nikola Jokic
big
Nikola Vucevic
big
Robin Lopez
big
Rudy Gobert
big
Serge Ibaka
big
Steven Adams
big
Taj Gibson
big
Thaddeus Young
scoring attempt play type
PRBallHandler
Isolation
OffScreenOrHandoff
Spotup
Postup
OffRebound
PRRollMan
Cut

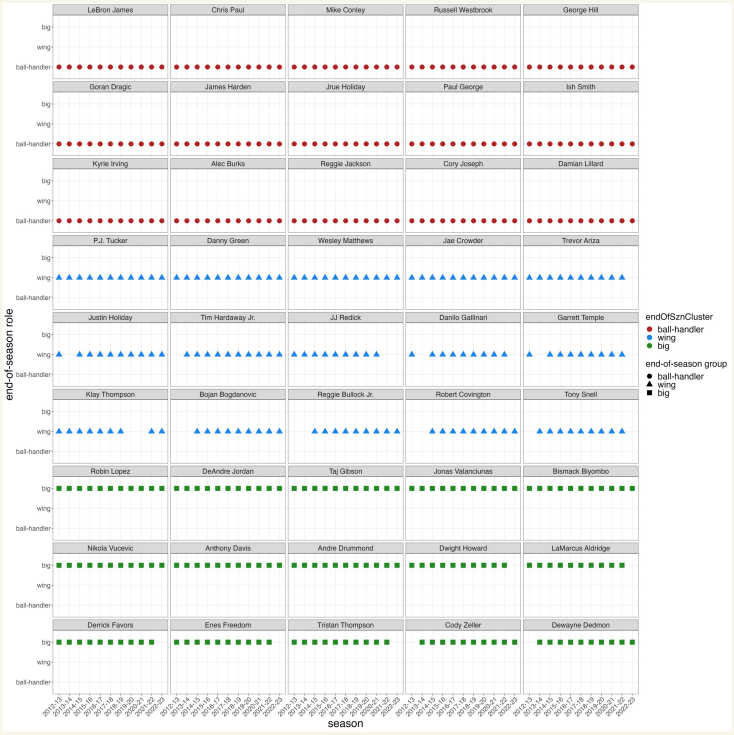
LeBron James
Chris Paul
Mike Conley
Russell Westbrook
George Hill
Goran Dragic
James Harden
Jrue Holiday
Paul George
Ish Smith
Kyrie Irving
Alec Burks
Reggie Jackson
Cory Joseph
Damian Lillard
P.J. Tucker
Danny Green
Wesley Matthews
Jae Crowder
Trevor Ariza
Justin Holiday
Tim Hardaway Jr.
JJ Redick
Danilo Gallinari
Garrett Temple
Klay Thompson
Bojan Bogdanovic
Reggie Bullock Jr.
Robert Covington
Tony Snell
Robin Lopez
DeAndre Jordan
Taj Gibson
Jonas Valanciunas
Bismack Biyombo
Nikola Vucevic
Anthony Davis
Andre Drummond
Dwight Howard
LaMarcus Aldridge
Derrick Favors
Enes Freedom
Tristan Thompson
Cody Zeller
Dewayne Dedmon
big
wing
ball-handler
end-of-season role
season
endOfSznCluster
ball-handler
wing
big
end-of-season group
ball-handler
wing
big

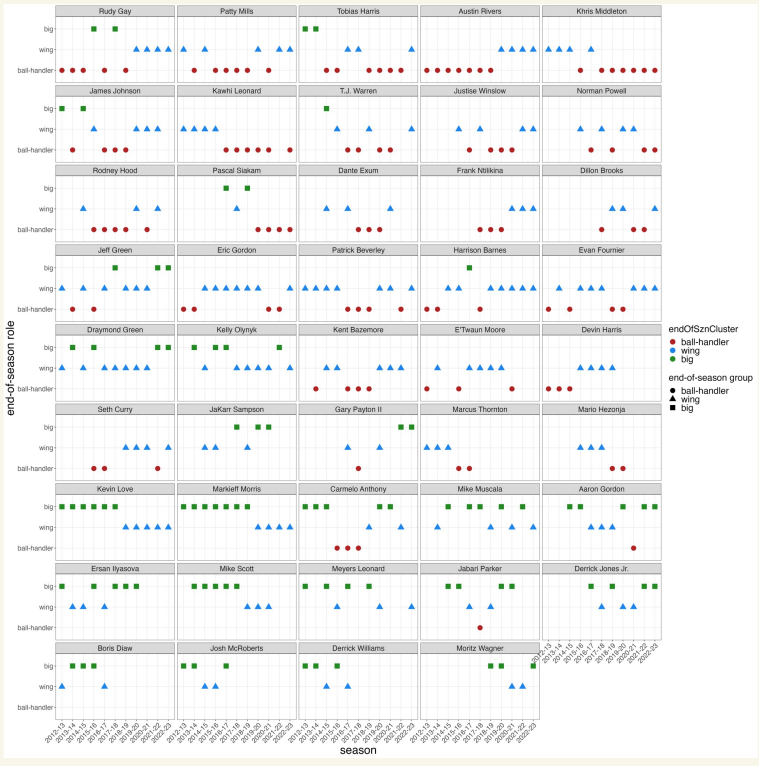
Rudy Gay
Patty Mills
Tobias Harris
Austin Rivers
Khris Middleton
James Johnson
Kawhi Leonard
T.J. Warren
Justise Winslow
Norman Powell
Rodney Hood
Pascal Siakam
Dante Exum
Frank Ntilikina
Dillon Brooks
Jeff Green
Eric Gordon
Patrick Beverley
Harrison Barnes
Evan Fournier
Draymond Green
Kelly Olynyk
Kent Bazemore
E'Twaun Moore
Devin Harris
Seth Curry
JaKarr Sampson
Gary Payton II
Marcus Thornton
Mario Hezonja
Kevin Love
Markieff Morris
Carmelo Anthony
Mike Muscala
Aaron Gordon
Ersan Ilyasova
Mike Scott
Meyers Leonard
Jabari Parker
Derrick Jones Jr.
Boris Diaw
Josh McRoberts
Derrick Williams
Moritz Wagner
big
wing
ball-handler
end-of-season role
season
endOfSznCluster
ball-handler
wing
big
end-of-season group
ball-handler
wing
big

I further clustered the 3 groups into more refined Roles: